Meng-Jiun Chiou

Taipei, Taiwan
✉ mengjiun.chiou@u.nus.edu
⚷ Personal Webpage
G Scholar ⚫ Github in Linkedin Ⓜ Medium

Summary

I am an applied scientist at Amazon Smart Home AI. I received my Ph.D. degree (Computer Science) from the National University of Singapore in 2022. My research interest during my Ph.D. is *learning structured representations of visual scenes* which covers visual relationship detection, scene graph generation and video understanding. I have 5 years+ experience in computer vision and machine learning research, while I am also experienced in Large Language Models and Supervised Fine-Tuning. I code in Python and PyTorch, and I am experienced in AWS cloud computing and MLOps.

Experience

Jul. 2022 – Present   **Applied Scientist**, *Amazon*, Taipei, Taiwan.
Researching and developing computer vision algorithms for intelligent devices at the Smart Home team under Amazon Devices and Services. Worked on i) home/room layout estimation via 3D reconstruction and 2D polygon fitting, ii) feasibility study of pet behavior understanding via video action detection, and iii) enhancing Alexa for a more ambient smart home experience with Large Language Models (LLM) and Supervised FineTuning (SFT).

Oct. 2020 – Jun. 2022   **Research Intern**, *TikTok*, Singapore.
Worked on i) unbiased scene graph generation via positive-unlabeled learning that achieves SOTA debiasing performance, ii) revealing the weakness of single-positive multi-label learning methods by adding real-world biases, and iii) improving smoking video detection by 10% at the Trust & Safety team.

Jun. 2020 – Oct. 2020   **Research Intern**, *ASUS Intelligent Cloud Services*, Singapore.
Worked on video human-object interaction (HOI) detection. Specifically, I introduced a new video HOI benchmark, *VidHOI* and proposed a spatial-temporal model *ST-HOI* which surpasses 2D/3D baselines.

Jul. 2013 – Jun. 2014   **Software Development Intern**, *Microsoft*, Taipei, Taiwan.
As a Microsoft Student Partner, I developed multiple Windows Apps, *e.g.*, *NHK Reader* with 7K+ downloads, and gave Microsoft Tech Talks on software development to Taiwan's college students.

Education

2017–2022   **Ph.D., Computer Science**, *National University of Singapore*, Singapore.
Supervised by Prof. Jiashi Feng and Prof. Roger Zimmermann. Worked on: Visual relationship detection, Scene graph generation, Human-object interaction recognition, Video (spatial-temporal) understanding

2012–2016   **B.Sc., Electrical and Computer Engineering**, *National Chiao Tung University (Currently, National Yang Ming Chiao Tung University)*, Hsinchu, Taiwan.
Overall GPA: 3.89/4.30 (or 3.90/4.00). Took various computer science courses.

2014–2015   **Exchange Program, Information & Communication Engineering**, *University of Tokyo*, Japan.
During the 1-year program, I worked on efficient look-up table based SVM classifiers for image classification at the *Multimedia Processing Lab*, supervised by Prof. Toshihiko Yamasaki and Prof. Kiyoharu Aizawa.

Publications

2022   **Meng-Jiun Chiou**. *Learning Structured Representations of Visual Scenes*. PhD thesis, National University of Singapore, 2022.

2021   **Meng-Jiun Chiou**, Roger Zimmermann, and Jiashi Feng. Visual relationship detection with visual-linguistic knowledge from multimodal representations. ***IEEE Access***, volume 9, pages 50441–50451. IEEE, 2021.

2021   **Meng-Jiun Chiou**, Chun-Yu Liao, Li-Wei Wang, Roger Zimmermann, and Jiashi Feng. St-hoi: A spatial-temporal baseline for human-object interaction detection in videos. In *Proceedings of the ACM International Conference on Multimedia Retrieval Workshops (**ACM ICMR-W'21**)*, pages 9–17, 2021.

2021   **Meng-Jiun Chiou**, Henghui Ding, Hanshu Yan, Changhu Wang, Roger Zimmermann, and Jiashi Feng. Recovering the unbiased scene graphs from the biased ones. In *Proceedings of the 29th ACM International Conference on Multimedia (**ACM MM'21**)*, pages 1581–1590, 2021.

2020 **Meng-Jiun Chiou**, Zhenguang Liu, Yifang Yin, An-An Liu, and Roger Zimmermann. Zero-shot multi-view indoor localization via graph location networks. In *Proceedings of the 28th ACM International Conference on Multimedia (**ACM MM'20**)*, pages 3431–3440, 2020.

2019 Yifang Yin, **Meng-Jiun Chiou**, Zhenguang Liu, Harsh Shrivastava, Rajiv Ratn Shah, and Roger Zimmermann. Multi-level fusion based class-aware attention model for weakly labeled audio tagging. In *Proceedings of the 27th ACM International Conference on Multimedia (**ACM MM'19**)*, pages 1304–1312, 2019.

2015 **Meng-Jiun Chiou**, Toshihiko Yamasaki, and Aizawa Kiyoharu. A fast table-based approach of bag-of-features for large-scale image classification. In *Proceedings of the ITE Annual Convention 2015 (**ITE'15**)*, pages 24A–1. The Institute of Image Information and Television Engineers, 2015.

2015 **Meng-Jiun Chiou**, Toshihiko Yamasaki, and Kiyoharu Aizawa. A fast method of visual words assignment of bag-of-features for object recognition. In *The 18th Meeting on Image Recognition and Understanding (**MIRU'15**)*, pages SS4–40, 2015.

Selected Projects

Affiliated with Amazon

2023 **Alexa Conversational Control for Smart Homes via Large Language Models**.
As an on-going project, we are working on an Alexa-related project aiming for a more ambient and conversational smart home experience with Large Language Models (LLM) and Supervised FineTuning (SFT). [Video] [Media Coverage by Bloomberg]

2023 **Alexa Map View: Room Layout Estimation via 3D Reconstruction and 2D Polygon Fitting**.
Worked on Alexa Map View (announced in Sep. 2023). Researched and prototyped a real-time home layout estimation system with RGB videos as input based on 3D mesh reconstruction methods, *e.g.*, NeuralRecon, and 2D polygon fitting approaches with rectilinear constraints for post-processing layouts, *e.g.* Floorplan Fitting. Worked on it end-to-end including literature review, algorithm implementation, metric definition and evaluation, and failure analysis. [Alexa Map View] [Video]

2022 **Pet Behavior Understanding by Action Classification & Spatio-Temporal Action Detection**.
We took an early-stage initiative to study pet behavior understanding. I worked on the whole end-to-end experiment pipeline, *i.e.*, problem definition (action classification/spatio-temporal action detection), data collection/cleaning/exploration, algorithm implementation, metric definition and error analysis. Training large 3D backbones with debiasing method like re-weighting, we obtained promising models for our goal.

Affiliated with TikTok (ByteDance AI Lab) & National University of Singapore

2022 **Improving Smoking Video Detection with new Architectures and Augmentations**.
We implemented SOTA data augmentation techniques including Mixup, Cutout and CutMix, and various new visual backbones such as Swin Transformer and we ended up *improving the smoking video detection performance by around 10 percent measured by recall*.

2021 **Revealing the biases in Single-Positive Multi-Label Learning**.
We revealed that the current Single-Positive Multi-Label (SPML) methods do not consider labeling bias such as *bounded rationality* and *reporting bias*, and we showed that *adding theses real-world biases to the existing SPML models would undermine their performance*. [Slides]

2021 **Unbiased Scene Graph Generation with Positive-Unlabeled Learning**.
We introduced *Dynamic Label Frequency Estimation* (DLFE) for debiasing scene graph generation (SGG). Applying DLFE to SGG methods we got *new SOTA debiasing performance*, i.e., *+5 averaged mean recall (24%→29%) or +21 tail-part recall (17%→38%)* v.s. previous SOTAs. [Paper] [Source Code] [Slides] [Poster] [Video]

Affiliated with ASUS Intelligent Cloud Services & National University of Singapore

2020 **Human-Object Interaction Detection in Videos**.
We *introduced* a keyframe-centered, large-scale video human-object interaction detection benchmark named *VidHOI*. Proposed a strong baseline called *ST-HOI outperforming the 2D/3D baseline models by obtaining 74% relatively or 6.1% absolutely higher mAP (8.3%→14.4%)* on temporal-related HOIs. [Paper] [Source Code & Dataset] [Slides] [Video]

Affiliated with National University of Singapore

2020 **Visual Relationship Detection with External Knowledge**.
We introduced a novel Transformer-based multi-modal visual relation detection architecture, named Relational Visual-Linguistic BERT (*RVL-BERT*), enriched by the visual-linguistic knowledge from large-scale external datasets. *RVL-BERT achieved SOTA performance* on the SpatialSense dataset and competitive results on the VRD and VG datasets. [Paper] [Source Code]

2019 ***Zero-Shot Indoor Localization with Floor Plans***.
We introduced a multi-view image-based indoor localization system named *GLN* achieving SOTA performance. Also proposed a zero-shot learning pipeline where we utilize the proposed *Map2Vec* location-aware embeddings. *Zero-shot GLN achieves promising results*, e.g., *56.3% 5-meter localization error.* [Paper] [Source Code] [Poster] [Video]

2018 ***Weakly-Labeled Audio Tagging with Attention-based Model***.
We introduced a multi-level attention-based audio tagging model making segment-level predictions with temporal modeling, followed by aggregations along both time and feature domains. *Our method achieves SOTA audio tagging results.* [Paper]

2017 ***Real-Time On-Device Blind Navigation***.
The Light navigates blind people to move around smoothly in real time using *MobileNet* for object segmentation. It *won 2nd prize at iNTUition Hackathon 2017.* [Project Page] [Source Code]

Affiliated with National Chiao Tung University

2016 ***Right Whale Idenditication with Fast R-CNN***.
We developed a right whale identification system face by training *Fast R-CNN* on a large-scale Kaggle dataset. [Technical Report] [Source Code]

Affiliated with Univeristy of Tokyo

2015 ***Fast Image Recognition with Look-Up Tabled-based Bag-of-Features***.
Table-Based Bag-of-Features (Table-Based BoF) is a fast look-up table based method for finding bag-of-features-based indexes of query pictures without feature extraction. [Paper] [Source Code]

Academic Services

2018–Present **Program Committee**, *NeurIPS Workshop on Distribution Shifts (DistShift) ('23/'22/'21), ACMMM'22 Open-Source Program, BigMM'20 Graduate Student Consortium, CVPR'18 Workshop on Visual Understanding of Humans in Crowd Scene*

2018–Present **Reviewer**, *ACMMM ('23/'22/'21/'20), IET Computer Vision ('23), IEEE TMM ('23/'21), Journal of Imaging ('23), IEEE TIP ('22/'20), ACM TOMM ('20), Springer MMSJ ('19)*, NUS MSCS Admission ('21/'20)

2017–2021 **Teaching Assistant**, Big-Data Analytics Technology (NUS, '21), Computer Vision and Pattern Recognition (NUS, '19/'18), Data Structures and Algorithms (NUS, '17), Special Friday Lecture for High School Students (UTokyo, '15)

Scholarships & Awards

2017 ***2nd Place***, *iNTUition Hackathon 2017*, a 24-hour hackathon at *Nanyang Technological University*.

2017 ***NUS Research Scholarship*** including full tuition waiver and monthly stipend, awarded by the *National University of Singapore*.

2015 ***Helm Technology Scholarship*** awarded by the *Helm Technology Inc.*, Taiwan.

2014 ***Student Exchange Support Program*** scholarship for exchange students to the *University of Tokyo*, awarded by *Japan Student Services Organization*.

2014 ***Short Term Exchange Scholarship*** for outbound exchange students, awarded by the *National Chiao Tung University*.

2014 ***Xiao Yuan-Long Scholarship*** for students with superb GPA, awarded by *National Chiao Tung University*.

Skills

Programming PyTorch, Python, Matlab, C, C++
Miscellaneous AWS Cloud Computing [Course Cert], MLOps [Course Cert], Docker
Language Mandarin Chinese (native speaker), English (fluent) and Japanese (fluent; JLPT N1)

Position of Responsibility

2013-2014 **Vice President**, *Chien-Kuo & Taipei First Girls' High School Alumni Association*, National Chiao Tung University.
I took on leadership roles to organize a variety of events for the two high schools' alumni.